

CPSC 331 - Assignment #3

- Let T be a binary search tree storing values of type V with keys of type E and let key be any value of type E . If the search algorithm is executed with inputs key and T , the algorithm eventually terminates and following are satisfied upon termination:
 - If key is a key stored in T then v , the value of the key we are searching for in T is returned.
 - If key is not stored in T then a `KeyNotFoundException` is thrown.
 - Neither T or key are changed

Thus, the given algorithm is correct for the search algorithm.

Loop Invariant:

- key is a value of type E
- T is a binary search tree with keys of type E and associated values of type V
- $curr \neq \text{null}$ and $v = \text{null}$
- key and T are unchanged
- The Binary Search Tree property is satisfied

Proof:

Suppose h , the height of a given tree T is -1 . Then T is an empty tree, so key cannot be stored within T . It is sufficient to show that the program halts with a `KeyNotFoundException` being thrown since the while loop is never entered and V is never re-assigned any value. T and key are never changed.

Since T is an empty tree, $curr$, which is initially set to the root of the tree T is set to `null`, since an empty tree has no root, and v is also set to `null`. So the loop test `while(curr != null and v = null)`, fails, so the loop body is skipped over, and we go to the if statement, `if(v = null)`, is true, so a `KeyNotFoundException` is thrown.

Now, suppose the height, h , is an integer greater than or equal to -1 , and suppose the search algorithm works correctly whenever its inputs consist of a key key of type E , and a tree T whose height is less than or equal to h . Let T be a binary search tree of height $h + 1$ greater than or equal to zero. We know T is not empty, so it contains some element with key or type E stored at the root of T (for which the key is accessible as $curr.key$, since $curr$ begins at the root of T and is not null since T is not empty).

Now, the key we are searching for is either less than $curr.key$, equal to $curr.key$, or greater than $curr.key$.

We consider all of these cases below:

Since the tree is not empty, $curr = T.root$, which is not null, and V is null, so the loop test passes.

- If key is equal to $curr.key$, then the following tests fail, so we assign the value of $curr.key$ to V , and the loop guard fails at the next iteration, so we skip over the loop body and go to the if statement, $if(v = null)$, which fails, since v is now equal to $curr.value$, so we execute the else, and return the value of v with key or T being changed.
- If key is less than $curr.key$, then we know that the first test failed, so $curr.key$ is not equal to key . So we assign the $curr.left$, the left child of $curr$, to $curr$. Now we must loop again, v is no longer null and:
 - Now key is either the key of a value stored in T or it is not.
 - If key is a value stored in T then it follows the binary search tree property and the key is stored somewhere in the left subtree of T , which is accessed by continually accessing $curr.left$. Since T has height $h+1$, the height of the left subtree is an integer between -1 and h , it now follows by the inductive hypothesis that the execution of the loop on inputs $curr = curr.left$ as the left subtree, $T.left$ and key terminates and $v = \text{value of } curr.key$ as the output without changing $curr$ and key as the output without changing key or the left subtree, which has been accessed by accessing $curr.left$.

The output produced, V , is immediately returned as the output of the algorithm, so it is clear that the algorithm terminates and returns an expected output without changing T or key .

- On the other hand, if *key* is not stored in *T* then it won't be stored in the left subtree, and we won't be able to access it by accessing *curr.left*. Since the height of the left subtree is an integer between -1 and *h*, it follows by the inductive hypothesis that the execution of the while on the left subtree by accessing *curr.left* and *v = null* terminates without assigning a value to *V*, and a *KeyNotFoundException* being thrown without changing the value of *T* or *key*.
- Finally, if *key* is greater than or equal to *curr.key*, then both of the previous tests "*curr.key = key*" and "*curr.key < key*" so we execute the else of the algorithm, so we assign the *curr.right*, the right child of *curr*, to *curr*. Now we must loop again, *v* is no longer null and either *key* is stored in *T* or it is not. Each subcase is considered:
 - If *key* in *T* then it follows by the Binary Search Tree Property that *key* is stored in the right subtree of *T*, which we may access by continually accessing *curr.right*. Since *T* has height *h + 1*, the height of the right subtree is an integer between -1 and *h*. It now follows by the inductive hypothesis that the execution of while loop on inputs *curr=curr.right* and *v = null* terminates and *v* is returned with the value of *curr.key* as output without changing *key* or *T*.

The produced output is returned as the output of the while loop and it is clear that the algorithm terminates and returned an expected output.

- On the other hand, if *key* is not stored in *T* then it won't be stored in the right subtree, and we won't be able to access it by accessing *curr.right*. Since the height of the right subtree is an integer between -1 and *h*, it follows by the inductive hypothesis that the execution of the while on the right subtree by accessing *curr.right* and *v = null* terminates without assigning a value to *V*, and a *KeyNotFoundException* being thrown without changing the value of *T* or *key*.

Thus the algorithm terminates and works correctly with inputs *T* and *key* in all cases, as is required.

Loop Variant: $f(h + 1, h) = h + 1 - h$.

- This function is integer valued
- The function decreases by at least one after every iteration, *h* increases so

- $(h+1) - h$ decreases.
- If $f(h+1, h)$ is less than or equal to zero, which is when $h = -1$, and the tree is empty, then the function terminates.

Thus, $f(h+1, h)$ terminates and its existence implies that the loop must terminate.

2.

insert(T, K key, V value)

curr = root;

parent = null;

if root = null

 root = new bstNode(key, value)

end if

else

 while(curr != null)

 parent = curr;

 if key < curr.key

 curr = curr.left;

 else if key > curr.key

 curr = curr.right

 else

 throw KeyFoundException

 end while

end else

if key < parent.key

 parent.left = bstNode(key, value);

else

 parent.right = bstNode(key, value);

3. inOrderTraversal(T)

if root = null

 end inorder traversal

Stack bstStack = new Stack

curr = T.root;

while stack.Empty = false or curr != null

 if curr != null

 bstStack.push(curr)

```

        curr = curr.left
    else
        node = bstStack.pop;
        curr = node.right;

```

7. The data obtained from running A3Q6.java is what I would expect to obtain.

- We know that a binary search tree's height is smallest when the tree is fullest, we can obtain this height using the equation: $\log_2(n+1)-1$.
- We also a binary search tree's height is the biggest when it is a continuous chain of length n, so the maximal height can be obtained from the equation $n-1$.
- If we use these equations to obtain the maximum and minimum height of 100 binary search trees of each n we tested we get:

	n=100	n=1000	n=10000	N=10000
Minimum	5.56	8.96	12.28	12.60
Maximum	99	999	9999	99999

Now, these values do differ quite a lot from the obtained values, however, one can expect that when n is sufficiently large and a binary search tree T is constructed with size n, it is very likely that the values of the maximum and the minimum will be much closer to the average height of m trees than the calculated maximum and minimum, since when constructing a tree randomly, it is extremely unlikely that we will get a completely full tree, or a continuous chain.

The upper-bound on the average is what I would expect, since it is calculated using the formula $3\log_2(n)$, and the calculated values are greater than that of the actual averages, so this value is as expected.

The worst-case bound on the maximum height of a random red-black tree of size n is also as expected, since the properties of a red-black tree to differ from those of a binary search tree, so the height will in fact differ, and this value was

calculated using the formula $\frac{\log_2(n+1)}{2}$